

# AI531 Final Project

By Mykyta Synytsia and Bach Xuan Phan




# Table of Contents

- Reversi (what it is and how it is played)
- How we are comparing the two adversarial agents
- Heuristics we employed
- Challenges we encountered and how we overcame them
- Results


# What is Reversi?

- A turn-based two-player game
- 8x8 grid
- Easy to learn but difficult to master
- Tree  $\sim 10^{60}$  nodes




# How to play Reversi?

- Initially, the four center cells are filled with four diagonally adjacent black and white discs
- Black goes first




# How to play Reversi?

- Players take turns placing discs




# How to play Reversi?

- For a disc to be placed legally, it must trap the opponent's discs vertically, horizontally or diagonally which results in those discs being flipped



# How to play Reversi?

- Once the grid is filled, or neither player can move, the player with the most discs wins.



# Fair comparison of MCTS and Minimax

To choose appropriate depths for minimax and iterations for MCTS we approximate the number of nodes explored for an 8x8 board per move

- Since the average branching factor is 10, a good approximation of the number of nodes explored for depth limited minimax would be  $10^L$ , where  $L$  is the depth
- In the case of MCTS,  $60 \cdot I$ , where  $I$  is the number of iterations



# Fair comparison of MCTS and Minimax

- From these approximations we chose depths and iterations that result in fairly equivalent node exploration numbers:
  - Depths of 3, 4, 5 and 6  $\Rightarrow 10^3, 10^4, 10^5$ , and  $10^6$
  - Iterations of 10, 100, 1000, 10000  $\Rightarrow 600, 6000, 60000, 600000$
- After every game, swap players

# Minimax evaluation function

$$score(n) = w_{pos} * \left( \sum n.discs(p) * w_b - \sum n.discs(o) * w_b \right) + w_{mob} * (n.moves(p) - n.moves(o))$$

- Positional Disc Score: Difference between the sum of the current player's discs and the opponent's discs
- Mobility Score: Difference in the number of possible moves between the current player and the opponent

$$score(n) = w_{pos} * \left( \sum n.discs(p) * w_b - \sum n.discs(o) * w_b \right) + w_{mob} * (n.moves(p) - n.moves(o))$$

- Where  $n$  is the current node and  $p$  and  $o$  correspond to player and opponent, respectively.
- $w_{pos}$  and  $w_{mob}$  correspond to a positional weight and mobility weight, we used a constant weight of 10 and 5.
- $n.discs(x)$  returns a matrix where indices are 1 for locations where player  $x$ 's disks are present, 0 otherwise.
- $W_b$  is a matrix with weights for each disc position.

$$score(n) = w_{pos} * \left( \sum n.discs(p) * w_b - \sum n.discs(o) * w_b \right) + w_{mob} * (n.moves(p) - n.moves(o))$$

- The number of valid moves in the current state for player x is designated as n.moves(x).
- With this heuristic evaluation function, scores are greater for the player when they have more discs on the board than the opponent and/or when the number of possible moves is greater.

# MCTS selection policy

$$UCB1(n) = \frac{wins(n)}{total\_games(n)} + c * \sqrt{\frac{\log(visits(n.parent))}{visits(n)}}$$

Exploitation

Exploration

$$UCB1(n) = \frac{wins(n)}{total\_games(n)} + C * \sqrt{\frac{\log(visits(n.parent))}{visits(n)}}$$

Exploitation

Exploration

- Upper confidence bound
- UCB1 selection policy for node n, balances exploration (exploring states which were visited less) and exploitation (exploring states which have performed well before).
- Helps to evaluate the utility of node n based on exploitation, win rate of node n, and exploration

$$UCB1(n) = \frac{wins(n)}{total\_games(n)} + C * \sqrt{\frac{\log(visits(n.parent))}{visits(n)}}$$

Exploitation

Exploration

- Along comes a balancing constant, C, a popular value for which is the square root of 2.
- The exploitation factor rises as the win rate increases, and when a node was not explored very much in relation to the parent, the exploration factor increases.
- This selection policy prioritizes resources to actions which statistically lead to better outcomes.

# Challenges: Time

- Single thread is impossible because it took too long.
- Python “threading” has Global Interpreter Lock, preventing multithread aside from I/O bound tasks => “multiprocessing”
- Simulations took a very long time. 1 preliminary test with just 8x8 took 48 hours. => Implement stop\_early



# Challenges: Multiprocessing Bug

- `AttributeError`: Python could not find `WEIGHTS` in `MinimaxPlayer`
- The child processes (used by `starmap`) won't inherit that `WEIGHTS` value
- But `WEIGHTS` is a static variable in `MinimaxPlayer`, that should mean it sticks through new processes

# Challenges: Multiprocessing Bug

Behavior	Linux/Unix (fork)	Windows (spawn)
<b>Memory sharing</b>	Child process inherits the <b>parent's memory</b> (Copy-on-Write).	Child starts <b>fresh</b> , no memory inheritance.

# Results

To evaluate the algorithms we focused on:

- Average total duration of searches
- Average number of nodes explored
- Win rate
- Searches that failed due to time limit

# Average duration of searches

## Minimax

Minimax average of total search time (s) vs board size and search depth				
Board Size	Depth			
	3	4	5	6
4	0.022	0.04	0.077	0.121
8	6.808	30.616	167.335	632.757
12	87.183	390.696	856.841	1376.846

## MCTS

MCTS average of total search time (s) vs board size and iterations				
Board Size	Iterations			
	10	100	1000	10000
4	0.044	0.299	1.834	12.024
8	4.062	39.075	370.572	1888.147
12	34.755	373.311	1772.112	2851.605

# Average of total nodes explored

## Minimax

Minimax average nodes explored vs board size and search depth				
Board Size	Depth			
	3	4	5	6
4	104.04	196	381.5	658.42
8	6924.04	32816.58	180567.2	745303
12	41911.46	189724.4	319085.7	666029

## MCTS

MCTS average nodes explored vs board size and iterations				
Board Size	Iterations			
	10	100	1000	10000
4	321.67	2237.54	12231.88	59256.71
8	8959.71	87142.33	836543.3	4195134
12	32708.46	361819.5	1598819	2729862

# Average minimax win rate

Minimax win rate vs board size						
Board Size	Depth	Iterations				Average
		10	100	1000	10000	
4	3	33.33%	33.33%	50.00%	50.00%	41.67%
	4	16.67%	0.00%	0.00%	0.00%	4.17%
	5	16.67%	0.00%	0.00%	0.00%	4.17%
	6	16.67%	0.00%	0.00%	0.00%	4.17%
8	3	100.00%	0.00%	16.67%	50.00%	41.67%
	4	100.00%	66.67%	16.67%	66.67%	62.50%
	5	100.00%	50.00%	16.67%	50.00%	54.17%
	6	100.00%	33.33%	66.67%	83.33%	70.83%
12	3	83.33%	66.67%	100.00%	100.00%	87.50%
	4	83.33%	33.33%	100.00%	83.33%	75.00%
	5	100.00%	100.00%	100.00%	83.33%	95.83%
	6	100.00%	83.33%	83.33%	66.67%	83.33%

# Failed searches (due to time limit)

Games stopped early vs board size, depth and MCTS iterations					
Board Size	Depth	Iterations			
		10	100	1000	10000
4	3	0	0	0	0
	4	0	0	0	0
	5	0	0	0	0
	6	0	0	0	0
8	3	0	0	0	6
	4	0	0	0	6
	5	0	0	0	6
	6	0	0	0	6
12	3	0	0	6	6
	4	0	0	5	6
	5	4	5	5	6
	6	6	5	6	6

# Other potential metrics for results

- Varying position/mobility weights for minimax evaluation function
- Different heuristic functions on minimax
- Varying balancing constant for UCB1



Thank you for  
your time!

